

The evolution of attention in AI

* Many real-world data types are sequences. In long-sequences, the important relationships can exist between elements far apart.

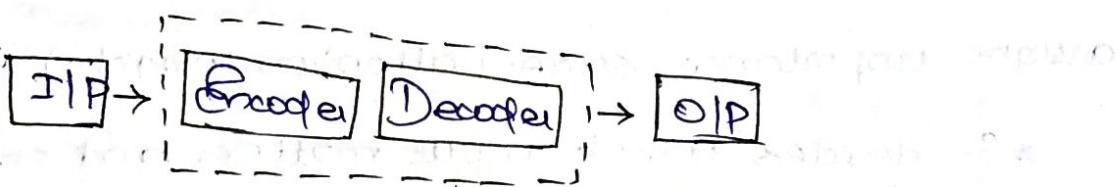
→ first and last word in a paragraph.

→ top-left and bottom-right patches in an image.

* Traditional models struggled with this problem.

Attention mechanisms were introduced to solve long-range dependency.

Seq2Seq model:



* A seq-2-seq model converts one IP sequence to another output sequence.

* Encoder is an RNN which processes the input sequentially, one at a time.

* The output of the encoder is a context vector. This single context vector represents the entire sentence meaning.

* The context vector is fed as an input to the decoder. The decoder generates one word at a time. It is also a RNN.

Problem in RNN Encoder-Decoder:

* Entire sentence meaning must fit in one vector, regardless of the sentence length.

* Decoder only sees the final context vector. It cannot see earlier hidden states / intermediate representation.

* This makes long-range dependencies extremely difficult.

Bahdanau Attention (2014):

* Instead of relying on one context vector, it allows the decoder to look at all encoder hidden states and assigns importance scores (attention weights).

* It decides which input matters and computes weighted contributions from all encoder states. This removes pressure from a single context vector.

* Bahdanau Attention = Attention b/w two sequences.

→ Used mainly in translation.

* Self Attention = Attention to itself.

- Every word looks at all other words.
- Produce context-aware embedding.